



Home

Study tools ▾

My courses ▾

My books

My folder

Career

Life



Find solutions for your homework

Search

home / study / engineering / computer science / computer science questions and answers / the bounded-buffer solution in the below code uses a...

Question: The bounded-buffer solution in the below code uses a last-in---

The bounded-buffer solution in the below code uses a last-in-first-out strategy (LIFO). Change the code to implement a FIFO (First-in-First-out) strategy. You may use the (in,out) pointer method (using semaphores to test if the queue is full or empty should alleviate the problem of only using up N-1 locations) or implement a FIFO queue. Use the correct counting semaphore implementation.

C code below

```
*****
*****

#include <stdio.h>
#include <stdlib.h>
#include <pthread.h>
#include <semaphore.h>

#define SIZE 5
#define NUMB_THREADS 6
#define PRODUCER_LOOPS 2

typedef int buffer_t;
buffer_t buffer[SIZE];
int buffer_index;

pthread_mutex_t buffer_mutex;
/* initially buffer will be empty. full_sem
will be initialized to buffer SIZE, which means
SIZE number of producer threads can write to it.
And empty_sem will be initialized to 0, so no
consumer can read from buffer until a producer
thread posts to empty_sem */
sem_t full_sem; /* when 0, buffer is full */
sem_t empty_sem; /* when 0, buffer is empty. Kind of
like an index for the buffer */

void insertbuffer(buffer_t value) {
    if (buffer_index < SIZE) {
        buffer[buffer_index++] = value;
    } else {
        printf("Buffer overflow\n");
    }
}

buffer_t dequeuebuffer() {
    if (buffer_index > 0) {
        return buffer[--buffer_index]; // buffer_index-- would be error!
    } else {
        printf("Buffer underflow\n");
    }
    return 0;
}

void *producer(void *thread_n) {
    int thread_numb = *(int *)thread_n;
    buffer_t value;
    int i=0;
    while (i++ < PRODUCER_LOOPS) {
        sleep(rand() % 10);
        value = rand() % 100;
        sem_wait(&full_sem); // sem=0: wait. sem>0: go and decrement it
        /* possible race condition here. After this thread wakes up,
        another thread could acquire mutex before this one, and add to list.
        Then the list would be full again
        and when this thread tried to insert to buffer there would be
        a buffer overflow error */
```

Post a question

Answers from our experts for your tough homework questions

[Continue to post](#)

20 questions remaining

My Textbook Solutions



Loose Leaf...

Principles of...

Essential...

16th Edition

2nd Edition

6th Edition

[View all solutions](#)



```
sem_post(&empty_sem); // post (increment) empty buffer semaphore
printf("Producer %d added %d to buffer\n", thread_numb, value);
}
pthread_exit(0);
}

void *consumer(void *thread_n) {
int thread_numb = *(int *)thread_n;
buffer_t value;
int i=0;
while (i++ < PRODUCER_LOOPS) {
sem_wait(&empty_sem);
/* there could be race condition here, that could cause
buffer underflow error */
pthread_mutex_lock(&buffer_mutex);
value = dequeuebuffer(value);
pthread_mutex_unlock(&buffer_mutex);
sem_post(&full_sem); // post (increment) full buffer semaphore
printf("Consumer %d dequeue %d from buffer\n", thread_numb, value);
}
pthread_exit(0);
}

int main(int argc, int **argv) {
buffer_index = 0;

pthread_mutex_init(&buffer_mutex, NULL);
sem_init(&full_sem, // sem_t *sem
0, // int pshared. 0 = shared between threads of process, 1 = shared between processes
SIZE); // unsigned int value. Initial value
sem_init(&empty_sem,
0,
0);
/* full_sem is initialized to buffer size because SIZE number of
producers can add one element to buffer each. They will wait
semaphore each time, which will decrement semaphore value.
empty_sem is initialized to 0, because buffer starts empty and
consumer cannot take any element from it. They will have to wait
until producer posts to that semaphore (increments semaphore
value) */
pthread_t thread[NUMB_THREADS];
int thread_numb[NUMB_THREADS];
int i;
for (i = 0; i < NUMB_THREADS; ) {
thread_numb[i] = i;
pthread_create(thread + i, // pthread_t *t
NULL, // const pthread_attr_t *attr
producer, // void *(*start_routine) (void *)
thread_numb + i); // void *arg
i++;
thread_numb[i] = i;
// playing a bit with thread and thread_numb pointers...
pthread_create(&thread[i], // pthread_t *t
NULL, // const pthread_attr_t *attr
consumer, // void *(*start_routine) (void *)
&thread_numb[i]); // void *arg
i++;
}

for (i = 0; i < NUMB_THREADS; i++)
pthread_join(thread[i], NULL);

pthread_mutex_destroy(&buffer_mutex);
sem_destroy(&full_sem);
sem_destroy(&empty_sem);

return 0;
}
```

Expert Answer ^①



Anonymous answered this
6,791 answers

Was this answer helpful?



Please find the answer which is as shown below, Thankyou.



```
#include <linux/pipe_fs_i.h>

static void wait_for_partner(struct inode* inode, unsigned int *cnt)
{
    int cur = *cnt;

    while (cur == *cnt) {
        pipe_wait(inode->i_pipe);
        if (signal_pending(current))
            break;
    }
}

static void wake_up_partner(struct inode* inode)
{
    wake_up_interruptible(&inode->i_pipe->wait);
}

static int fifo_open(struct inode *inode, struct file *filp)
{
    struct pipe_inode_info *pipe;
    int ret;

    mutex_lock(&inode->i_mutex);
    pipe = inode->i_pipe;
    if (!pipe) {
        ret = -ENOMEM;
        pipe = alloc_pipe_info(inode);
        if (!pipe)
            goto err_nocleanup;
        inode->i_pipe = pipe;
    }
    filp->f_version = 0;

    /* We can only do regular read/write on fifos */
    filp->f_mode &= (FMODE_READ | FMODE_WRITE);

    switch (filp->f_mode) {
    case FMODE_READ:
        /*
         * O_RDONLY
         * POSIX.1 says that O_NONBLOCK means return with the FIFO
         * opened, even when there is no process writing the FIFO.
         */
        filp->f_op = &read_pipefifo_fops;
        pipe->r_counter++;
        if (pipe->readers++ == 0)
            wake_up_partner(inode);

        if (!pipe->writers) {
            if ((filp->f_flags & O_NONBLOCK)) {
                /* suppress POLLHUP until we have
                 * seen a writer */
                filp->f_version = pipe->w_counter;
            } else {
                {
                    wait_for_partner(inode, &pipe->w_counter);
                    if (signal_pending(current))
                        goto err_rd;
                }
            }
        }
        break;

    case FMODE_WRITE:
        /*
         * O_WRONLY
         * POSIX.1 says that O_NONBLOCK means return -1 with
         * errno=ENXIO when there is no process reading the FIFO.
         */
        ret = -ENXIO;
        if ((filp->f_flags & O_NONBLOCK) && !pipe->readers)
            goto err;

        filp->f_op = &write_pipefifo_fops;
        pipe->w_counter++;
        if (!pipe->writers++)
            wake_up_partner(inode);
    }
```



```

        if (signal_pending(current))
            goto err_wr;
    }
    break;

case FMODE_READ | FMODE_WRITE:
/*
 * O_RDWR
 * POSIX.1 leaves this case "undefined" when O_NONBLOCK is set.
 * This implementation will NEVER block on a O_RDWR open, since
 * the process can at least talk to itself.
 */
    filp->f_op = &rdwr_pipefifo_fops;

    pipe->readers++;
    pipe->writers++;
    pipe->r_counter++;
    pipe->w_counter++;
    if (pipe->readers == 1 || pipe->writers == 1)
        wake_up_partner(inode);
    break;

default:
    ret = -EINVAL;
    goto err;
}

/* Ok! */
mutex_unlock(&inode->i_mutex);
return 0;

err_rd:
if (!--pipe->readers)
    wake_up_interruptible(&pipe->wait);
ret = -ERESTARTSYS;
goto err;

err_wr:
if (!--pipe->writers)
    wake_up_interruptible(&pipe->wait);
ret = -ERESTARTSYS;
goto err;

err:
if (!pipe->readers && !pipe->writers)
    free_pipe_info(inode);

err_nocleanup:
mutex_unlock(&inode->i_mutex);
return ret;
}

/*
 * Dummy default file-operations: the only thing this does
 * is contain the open that then fills in the correct operations
 * depending on the access mode of the file...
 */
const struct file_operations def_fifo_fops = {
    .open      = fifo_open, /* will set read_ or write_pipefifo_fops */
};

```

Hope you got some explanation, still having doubts, mention in comment section, happy to resolve them,
Thankyou.

Thankyou. Happy Learning.

[View comments \(1\)](#) >

Questions viewed by other students

Q: A condensate line 152 mm nominal size made of schedule 40 carbon steel pipe is supported by threaded rod hangers spaced at 2.5 m center-to-center. The hangers are carbon steel, 50 cm long, with a root diameter of 12 mm.

Q: Problem 2. If A_n is a set for each $n \in \mathbb{N}$, then $\bigcap_{n \in \mathbb{N}} A_n$ denotes the set $\{x : \forall n \in \mathbb{N} \ x \in A_n\}$. (This is a special case of a more general "indexed intersection", where the \mathbb{N} can be replaced by any "index set"; see Section 1.8 of the textbook for more details and examples). Prove that $\bigcap_{n \in \mathbb{N}} [3, 3+n] = [3, 3]$. In other words, prove that $\bigcap_{n \in \mathbb{N}} A_n$ is the half-open interval $[3, 3]$.

A: [See answer](#) 100% (1 rating)

[Show more](#)

COMPANY

LEGAL & POLICIES

CHEGG PRODUCTS AND SERVICES

CHEGG NETWORK

CUSTOMER SERVICE



© 2003-2022 Chegg Inc. All rights reserved.